

Mapping Techniques for Load Balancing



Parallelization

- Decomposition
- Mapped to processes
- Overhead reduction

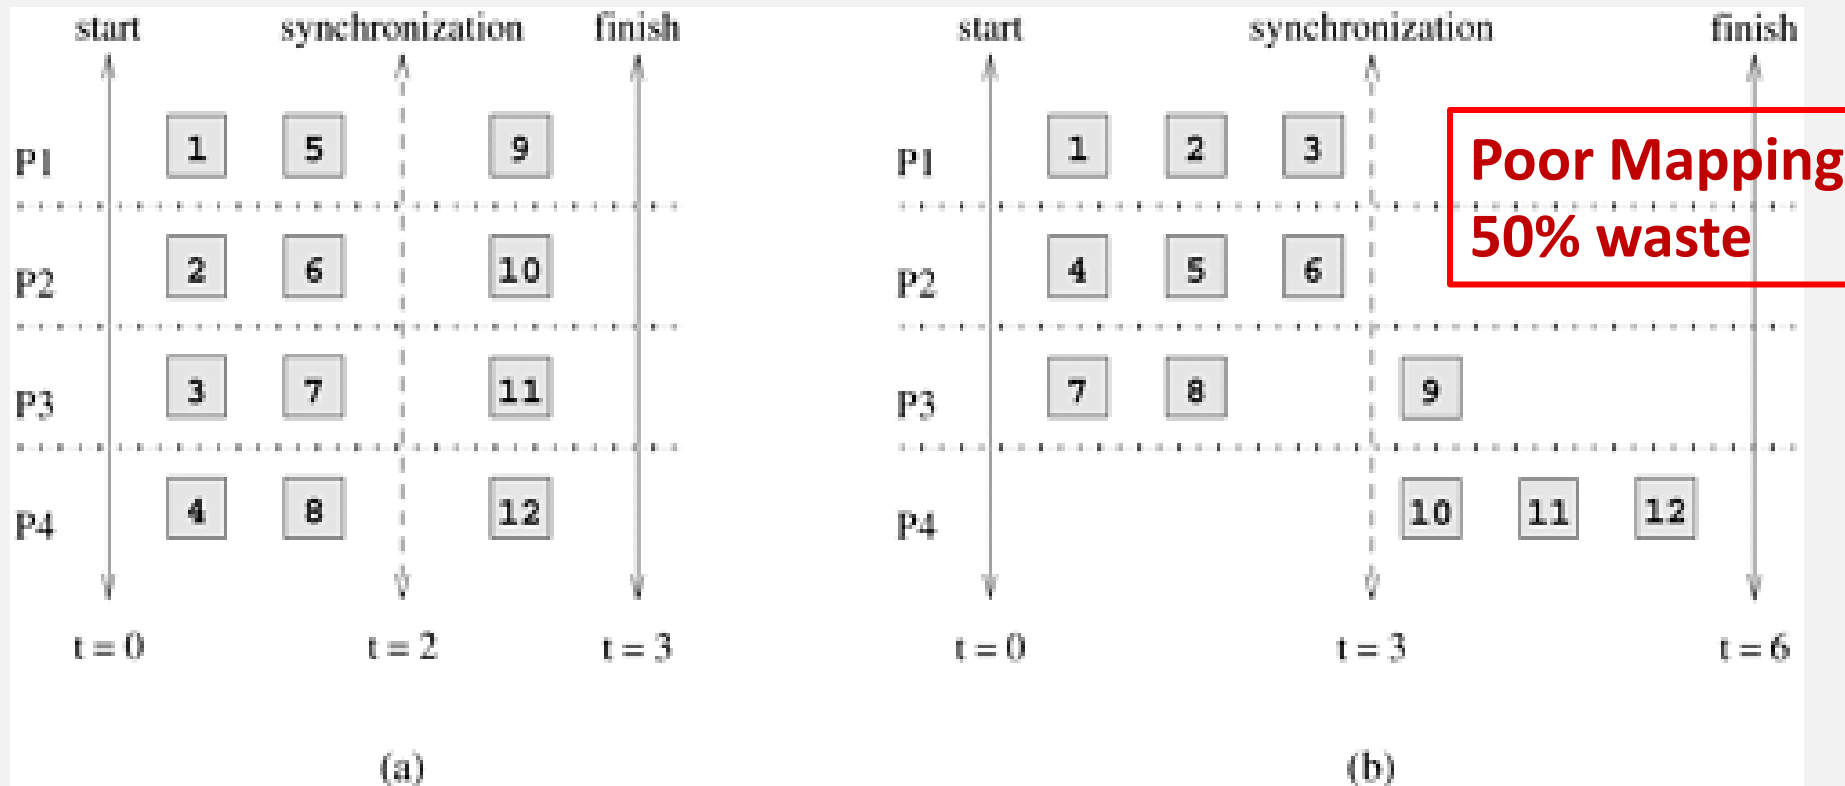
Good Mapping

- Finding a good mapping is a nontrivial problem
- Twin objectives
 - reducing the amount of time processes spend in interacting with each other,
 - reducing the total amount of time some processes are idle while the others are engaged in performing some tasks.
- They conflict each other
 - Less interaction---->unbalanced workload
 - Good balance--→ high interaction

Good Mapping

- A good mapping must ensure that the **computations** and **interactions** among **processes** at each **stage** of the **execution** of the parallel algorithm are **well balanced**

Two mappings of a hypothetical decomposition with a synchronization.



Task 9,10,11,12 can start after completion of first eight tasks
Different completion times for the two mappings

Two categories

- Static
- Dynamic

Static

- Used in conjunction with a decomposition based on **data partitioning**
- Used for mapping certain problems that are expressed naturally by a **static task-dependency graph**

Mappings Based on Data Partitioning

- Mappings based on partitioning two of the most common ways of representing data in algorithms, namely, **arrays and graphs**

Array Distribution Schemes

- The tasks are closely associated with portions of data by the **owner-computes rule**
- Therefore, mapping the relevant data onto the processes is equivalent to mapping tasks onto processes

Block Distributions

- A d -dimensional array is distributed among the processes
 - such that each process receives a contiguous block of array entries along a specified subset of array dimensions.
- Block distributions of arrays are particularly suitable when there is a locality of interaction, i.e.,
 - computation of an element of an array requires other nearby elements in the array.

1D partitioning among 8 processes

row-wise distribution

P_0
P_1
P_2
P_3
P_4
P_5
P_6
P_7


column-wise distribution

P_0	P_1	P_2	P_3	P_4	P_5	P_6	P_7
-------	-------	-------	-------	-------	-------	-------	-------


For example, consider an $n \times n$ two-dimensional array A with n rows and n columns. Now select one of these dimensions, e.g., the first dimension, and partition the array into p parts such that the k th part contains rows $kn/p \dots (k+1)n/p - 1$, where $0 \leq k < p$. That is, each partition contains a block of n/p consecutive rows of A . Similarly, if we partition A along the second dimension, then each partition contains a block of n/p consecutive columns.

Block Distribution and Data Sharing for Dense Matrix Multiplication


- Row-based 1-D



- Column-based 1-D



- Row/Col-based 2-D



P_0	P_1	P_2	P_3
P_4	P_5	P_6	P_7
P_8	P_9	P_{10}	P_{11}
P_{12}	P_{13}	P_{14}	P_{15}

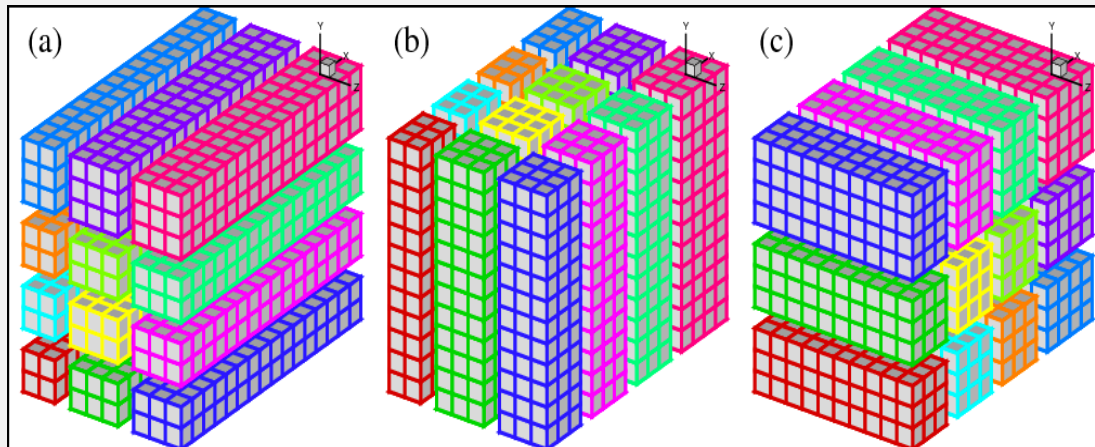
(a)

P_0	P_1	P_2	P_3	P_4	P_5	P_6	P_7
P_8	P_9	P_{10}	P_{11}	P_{12}	P_{13}	P_{14}	P_{15}

(b)

Two-dimensional distributions of an array, (a) on a 4 x 4 process grid, and (b) on a 2 x 8 process grid.


Instead of selecting a single dimension, we can **select multiple dimensions** to partition. For instance, in the case of array **A** we can select both dimensions and partition the matrix into blocks such that each block corresponds to a $n/p_1 \times n/p_2$ section of the matrix, with $p = p_1 \times p_2$ being the number of processes




- For example, consider the $n \times n$ matrix multiplication $C = A \times B$
- One way of decomposing this computation is to partition the output matrix C .
- Since each entry of C requires the same amount of computation, we can balance the computations by using either a one- or two-dimensional block distribution to partition C uniformly among the p available processes.
- In the first case, each process will get a block of n/p rows (or columns) of C , whereas in the second case, each process will get a block of size $\frac{n}{\sqrt{p}} \times \frac{n}{\sqrt{p}}$.
- In either case, the process will be responsible for computing the entries of the partition of C assigned to it.

- In the case of matrix-matrix multiplication, a one-dimensional distribution will allow us to use up to n processes by assigning a single row of C to each process.
- On the other hand, a two-dimensional distribution will allow us to use up to n^2 processes by assigning a single element of C to each process.

- In addition to allowing a higher degree of concurrency, **higher dimensional distributions** also sometimes help in **reducing the amount of interactions** among the different processes for many problems.



(a)



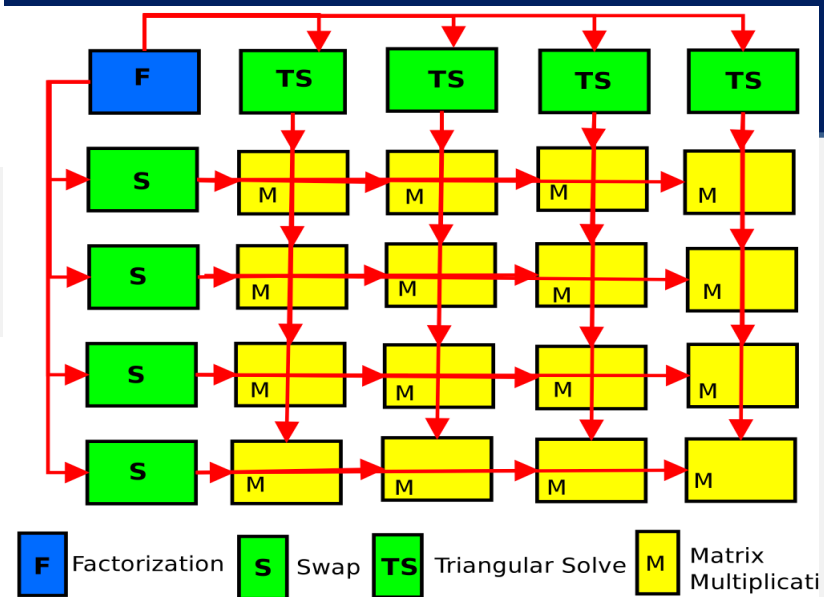

(b)

- With a one-dimensional partitioning along the rows, each process needs to access the corresponding n/p rows of matrix A and the entire matrix B , for process P_5 .
- Thus the total amount of data that needs to be accessed is $n^2/p + n^2$

- However, with a two-dimensional distribution, each process needs to access $\frac{n}{\sqrt{p}}$ rows of matrix A and $\frac{n}{\sqrt{p}}$ columns of matrix B for process P_5 .
- In the two-dimensional case, the total amount of shared data that each process needs to access is $\frac{n^2}{\sqrt{p}}$, which is significantly smaller compared to $O(n^2)$ shared data in the one-dimensional case.

Cyclic and Block Cyclic Distributions

- If the **amount of work differs** for different elements of a matrix, a **block distribution** can potentially lead to **load imbalances**.
- A classic example of this phenomenon is **LU factorization** of a matrix, in which the amount of **computation increases** from the **top left to the bottom right** of the matrix.



```

1.  procedure COL_LU (A)
2.  begin
3.      for k := 1 to n do
4.          for j := k to n do
5.               $A[j, k] := A[j, k] / A[k, k];$ 
6.          endfor;
7.          for j := k + 1 to n do
8.              for i := k + 1 to n do
9.                   $A[i, j] := A[i, j] - A[i, k] \times A[k, j];$ 
10.             endfor;
11.          endfor;

```

/*

After this iteration, column $A[k + 1 : n, k]$ is logically the k th column of L and row $A[k, k : n]$ is logically the k th row of U .

*/

LU Factorization of a Dense Matrix

A decomposition of LU factorization into 14 tasks

$$\begin{pmatrix} A_{1,1} & A_{1,2} & A_{1,3} \\ A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,1} & A_{3,2} & A_{3,3} \end{pmatrix} \rightarrow \begin{pmatrix} L_{1,1} & 0 & 0 \\ L_{2,1} & L_{2,2} & 0 \\ L_{3,1} & L_{3,2} & L_{3,3} \end{pmatrix} \cdot \begin{pmatrix} U_{1,1} & U_{1,2} & U_{1,3} \\ 0 & U_{2,2} & U_{2,3} \\ 0 & 0 & U_{3,3} \end{pmatrix}$$

1: $A_{1,1} \rightarrow L_{1,1}U_{1,1}$

2: $L_{2,1} = A_{2,1}U_{1,1}^{-1}$

3: $L_{3,1} = A_{3,1}U_{1,1}^{-1}$

4: $U_{1,2} = L_{1,1}^{-1}A_{1,2}$

5: $U_{1,3} = L_{1,1}^{-1}A_{1,3}$

6: $A_{2,2} = A_{2,2} - L_{2,1}U_{1,2}$

7: $A_{3,2} = A_{3,2} - L_{3,1}U_{1,2}$

8: $A_{2,3} = A_{2,3} - L_{2,1}U_{1,3}$

9: $A_{3,3} = A_{3,3} - L_{3,1}U_{1,3}$

10: $A_{2,2} \rightarrow L_{2,2}U_{2,2}$

11: $L_{3,2} = A_{3,2}U_{2,2}^{-1}$

12: $U_{2,3} = L_{2,2}^{-1}A_{2,3}$

13: $A_{3,3} = A_{3,3} - L_{3,2}U_{2,3}$

14: $A_{3,3} \rightarrow L_{3,3}U_{3,3}$

Block Distribution for LU

Notice the significant load imbalance

P₀ T ₁	P₃ T ₄	P₆ T ₅
P₁ T ₂	P₄ T ₆ T ₁₀	P₇ T ₈ T ₁₂
P₂ T ₃	P₅ T ₇ T ₁₁	P₈ T ₉ T ₁₃ T ₁₄

Block Cyclic Distributions



- Variation of the block distribution scheme
 - Partition an array into many more blocks (i.e. tasks) than the number of available processes.
 - Blocks are assigned to processes in a round-robin manner so that each process gets several non-adjacent blocks.
 - $N-1$ mapping of tasks to processes

Block-Cyclic Distribution for Gaussian Elimination

- Active submatrix shrinks as elimination progresses
- Assigning blocks in a block-cyclic fashion
 - Each PEs receives blocks from different parts of the matrix
 - In one batch of mapping, the PE doing the most will most likely receive the least in the next batch


Block-Cyclic Distribution

- A cyclic distribution: a special case with block size = 1
- A block distribution: a special case with block size = n/p
- n is the dimension of the matrix and p is the #of processes



Block Partitioning and Random Mapping

- Sparse matrix computations



(a)

P_0	P_1	P_2	P_3	P_0	P_1	P_2	P_3
P_4	P_5	P_6	P_7	P_4	P_5	P_6	P_7
P_8	P_9	P_{10}	P_{11}	P_8	P_9	P_{10}	P_{11}
P_{12}	P_{13}	P_{14}	P_{15}	P_{12}	P_{13}	P_{14}	P_{15}
P_0	P_1	P_2	P_3	P_0	P_1	P_2	P_3
P_4	P_5	P_6	P_7	P_4	P_5	P_6	P_7
P_8	P_9	P_{10}	P_{11}	P_8	P_9	P_{10}	P_{11}
P_{12}	P_{13}	P_{14}	P_{15}	P_{12}	P_{13}	P_{14}	P_{15}

(b)


Load imbalance using block-cyclic partitioning/mapping
– more **non-zero blocks** to diagonal processes **P_0 , P_5 , P_{10} , and P_{15}** than others
– P_{12} gets nothing

Block Partitioning and Random Mapping


$V = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]$

$\text{random}(V) = [8, 2, 6, 0, 3, 7, 11, 1, 9, 5, 4, 10]$

mapping = 8 2 6 0 3 7 11 1 9 5 4 10
 └─┬─┘└─┬─┘└─┬─┘└─┬─┘
 P₀ P₁ P₂ P₃



(a)



(b)

P_0	P_1	P_2	P_3
P_4	P_5	P_6	P_7
P_8	P_9	P_{10}	P_{11}
P_{12}	P_{13}	P_{14}	P_{15}

(c)

Graph Partitioning Based Data Decomposition


- Array-based partitioning and static mapping
 - Regular domain, i.e. rectangular, mostly dense matrix
 - Structured and regular interaction patterns
 - Quite effective in balancing the computations and minimizing the interactions
- Irregular domain
 - Sparse matrix-related
 - Numerical simulations of physical phenomena
 - Car, water/blood flow, geographic
- Partition the irregular domain so as to
 - Assign equal number of nodes to each process
 - Minimizing edge count of the partition.

Schemes for Dynamic Mapping

- Also referred to as dynamic load balancing
- Load balancing is the primary motivation for dynamic mapping.
- Dynamic mapping schemes can be
 - Centralized
 - Distributed

Centralized Dynamic Mapping

- Processes are designated as **masters** or **slaves**
 - Workers (slave is politically incorrect)
- General strategies
 - Master has pool of tasks and as central dispatcher
 - When one runs out of work, it requests from master for more work.
- Challenge
 - When process # increases, master may become the bottleneck.
- Approach
 - Chunk scheduling: a process picks up multiple tasks at once
 - Chunk size:
 - Large chunk sizes may lead to significant load imbalances as well
 - Schemes to gradually decrease chunk size as the computation progresses.



Distributed Dynamic Mapping

- All processes are created equal
 - Each can send or receive work from others
 - Alleviates the bottleneck in centralized schemes.
- Four critical design questions:
 - how are sending and receiving processes paired together
 - who initiates work transfer
 - how much work is transferred
 - when is a transfer triggered?

Reference

- <http://parallelcomp.uw.hu/ch03lev1sec4.html>
- https://passlab.github.io/CSCE569/notes/lecture10_design02.pdf