

# Module 4: Communication




# Distributed Shared Memory (DSM) Approaches

# How to share data between distributed processor nodes?

- Distributed Shared Memory
  - It's the abstraction to share data between distributed computers
- It enables access across multiple distributed physical memory to be accessed as single shared memory

# DSM Abstraction



# Working of DSM

- DSM enables easy access of individual computer shared data items
- DSM run time support sends updates as messages between computers
- Each computer has a local copy of recently accessed data items stored in DSM

# ...Continued...

- Shared memory request for a non-local piece of data is raised
- Single copy of data fetched and given to the requested system
- If multiple machines access the data at the same time, synchronization primitive like semaphore is used to handle the situation

# ...Continued

- Read(shared variable)
- Write (data, shared variable)

# Issues related to DSM Semantics

- Structure and granularity – data shared at the bit, word or page level
- Consistency – If multiple requests for a single data and each machine tries to update it, consistency should be maintained. This involves cache coherence like solution
- Heterogeneity – Accommodating different data representations of different machines, languages and OS
- Scalability – Bus latency, Increased broadcast messages



# DSM Versus Message Passing

- No marshaling of messages in DSM where as messages are marshalled and unmarshalled in message passing
- Synchronization in DSM is by constructs like semaphore/locks where as in message passing it is by message passing primitives
- DSM → processes can communicate with non-overlapping lifetimes where as in message passing processes communicate at the same time
- Efficiency in DSM heavily depends on the pattern of Data access by multiple machines where as the same is not true for message passing

# DSM features

- Space- uncoupled
- Time – uncoupled
- State based service
- Used for parallel and distributed computation
- Limited scalability
- Not associative

# Distributed Resource Management: Distributed Shared Memory

Courtesy: CS-550: Distributed Shared  
Memory [SiS '94]


# DSM

- What
  - The distributed shared memory (DSM) implements the shared memory model in distributed systems, which have no physical shared memory
  - The shared memory model provides a virtual address space shared between all nodes
  - To overcome the high cost of communication in distributed systems, DSM systems move data to the location of access

# DSM

- How:
  - Data moves between main memory and secondary memory (within a node) and between main memories of different nodes
  - Each data object is owned by a node
    - Initial owner is the node that created object
    - Ownership can change as object moves from node to node
  - When a process accesses data in the shared address space, the mapping manager maps shared memory address to physical memory (local or remote)

# DSM



# Advantages of distributed shared memory (DSM)

- Data sharing is **implicit**, hiding data movement (as opposed to 'Send'/'Receive' in message passing model)
- **Passing data structures** containing pointers is **easier** (in message passing model data moves between different address spaces)
- Moving entire object to user takes **advantage of locality difference**
- **Less expensive** to build than tightly coupled multiprocessor system: off-the-shelf hardware, no expensive interface to shared physical memory
- Very **large total physical memory** for all nodes: Large programs can run more efficiently
- **No serial access to common bus** for shared physical memory like in multiprocessor systems
- **Programs written for shared memory** multiprocessors can be **run on DSM** systems with **minimum changes**

# Algorithms for implementing DSM

- Issues
  - How to keep **track** of the location of remote data
  - How to minimize **communication overhead** when accessing remote data
  - How to access **concurrently** remote data at several nodes



# The Central Server Algorithm

- Central server maintains all shared data
  - **Read request**: returns data item
  - **Write request**: updates data and returns acknowledgement message
- Implementation
  - A **timeout** is used to resend a request if **acknowledgment fails**
  - Associated sequence numbers can be used to **detect duplicate write requests**
  - If an **application's request** to access shared data **fails repeatedly**, a **failure condition is sent** to the application
- Issues: performance and reliability
- Possible solutions
  - **Partition** shared data between several servers
  - Use a **mapping function** to distribute/locate data

# The Migration Algorithm

- Operation
  - **Ship** (migrate) entire data object (page, block) containing **data** item to requesting location
  - Allow **only one node** to access a shared data at a time
- Advantages
  - Takes advantage of the **locality of reference**
  - DSM can be **integrated with VM** at each node
    - Make DSM page multiple of VM page size
    - A locally held shared memory can be mapped into the VM page address space
    - If page not local, **fault-handler** migrates page and removes it from address space at remote node

# Migration Algorithm

- To locate a remote data object:
  - Use a location server
  - Maintain hints at each node
  - Broadcast query
- Issues
  - Only one node can access a data object at a time
  - **Thrashing** can occur: to minimize it, set minimum time data object resides at a node

# The Read-Replication Algorithm

- Replicates data objects to multiple nodes
- DSM keeps **track of location** of data objects
- Multiple nodes can have read access or one node write access (**multiple readers-one writer protocol**)
- After a write, **all copies are invalidated or updated**
- DSM has to keep track of locations of all copies of data objects. Examples of implementations:
  - **IVY**: owner node of data object knows all nodes that have copies
  - **PLUS**: distributed linked-list tracks all nodes that have copies
- Advantage
  - The read-replication can lead to substantial performance improvements if the ratio of reads to writes is large

# The Full-Replication Algorithm

- Extension of read-replication algorithm: multiple nodes can read and multiple nodes can write (**multiple-readers, multiple-writers protocol**)
- Issue: consistency of data for multiple writers
- Solution: use of **gap-free sequencer**
  - All writes sent to sequencer
  - **Sequencer** assigns sequence number and sends write request to all sites that have copies
  - Each node performs **writes** according to **sequence numbers**
  - A **gap** in sequence numbers indicates a missing write request: node asks for **retransmission** of missing write requests

# Memory coherence

- DSM are based on
  - Replicated shared data objects
  - Concurrent access of data objects at many nodes
- Coherent memory: when value returned by read operation is the expected value (e.g., value of most recent write)
- Mechanism that control/synchronizes accesses is needed to maintain memory coherence
- Sequential consistency: A system is sequentially consistent if
  - The result of any execution of operations of all processors is the same as if they were executed in sequential order, and
  - The operations of each processor appear in this sequence in the order specified by its program
- General consistency:
  - All copies of a memory location (replicas) eventually contain same data when all writes issued by every processor have completed

# Memory coherence (Cont.)

- Processor consistency:
  - Operations issued by a processor are performed in the order they are issued
  - Operations issued by several processors may not be performed in the same order (e.g. simultaneous reads of same location by different processors may yields different results)
- Weak consistency:
  - Memory is consistent only (immediately) after a synchronization operation
  - A regular data access can be performed only after all previous synchronization accesses have completed

# Memory coherence (Cont.)

- Release consistency:
  - Further relaxation of weak consistency
  - Synchronization operations must be consistent with each other only within a processor
  - Synchronization operations: Acquire (i.e. lock), Release (i.e. unlock)
  - Sequence:       Acquire  
                      Regular access  
                      Release



# Coherence Protocols

- Issues
  - How do we ensure that all replicas have the same information
  - How do we ensure that nodes do not access stale data

# Write-invalidate protocol

- A write to shared data invalidates all copies except one before write executes
- Invalidated copies are no longer accessible
- Advantage: good performance for
  - Many updates between reads
  - Per node locality of reference
- Disadvantage
  - Invalidations sent to all nodes that have copies
  - Inefficient if many nodes access same object
- Examples: most DSM systems: IVY, Clouds, Dash, Memnet, Mermaid, and Mirage

# Write-update protocol

- A write to shared data causes all copies to be updated (new value sent, instead of validation)
- More difficult to implement

# Design issues

- Granularity: size of shared memory unit
  - If DSM page size is a multiple of the local virtual memory (VM) management page size (supported by hardware), then DSM can be integrated with VM, i.e. use the VM page handling
  - Advantages vs. disadvantages of using a large page size:
    - (+) Exploit locality of reference
    - (+) Less overhead in page transport
    - (-) More contention for page by many processes
  - Advantages vs. disadvantages of using a small page size
    - (+) Less contention
    - (+) Less false sharing (page contains two items, not shared but needed by two processes)
    - (-) More page traffic
  - Examples
    - PLUS: page size 4 Kbytes, unit of memory access is 32-bit word
    - Clouds, Munin: object is unit of shared data structure

# Design issues (cont.)

- Page replacement
  - Replacement algorithm (e.g. LRU) must take into account page access modes: shared, private, read-only, writable
  - Example: LRU with access modes
    - Private (local) pages to be replaced before shared ones
    - Private pages swapped to disk
    - Shared pages sent over network to owner
    - Read-only pages may be discarded (owners have a copy)

# Case studies: IVY

- IVY (Integrated shared Virtual memory at Yale) implemented in Apollo DOMAIN environment, i.e. Apollo workstations on a token ring
- Granularity: 1 Kbyte page
- Process address space: private space + shared VM space
  - Private space: local to process
  - Shared space: can be accessed by any process through the shared part of its address space

# Case studies: IVY

- Node mapping manager: does mapping between local memory of that node and the shared virtual memory space
- Memory access operation
  - On page fault, block process
  - If page local, fetch from secondary memory
  - If not local, request a remote memory access, acquire page
- Page now available to all processes at the node

# Case studies: IVY (Cont.)

- Coherence protocol
  - Page access modes: read only, write, nil (invalidate)
  - Multiple readers-single writer semantics
  - Protocol
    - Write invalidation: before a write to a page is allowed, all other read-only copies are invalidated
    - Strict consistency: a reader always sees the latest value written



# Case studies: IVY (Cont.)

- Write sequence
  - Processor 'i' has write fault to page 'p'
  - Processor 'i' finds owner of page 'p' and sends request
  - Owner of 'p' sends page and its copyset to 'i' and marks 'p' entry in its page table 'nil' (copyset = list of processors containing read-only copy of page)
  - Processor 'i' sends invalidation messages to all processors in copyset

# Case studies: IVY (Cont.)

- Read sequence
  - Processor 'i' has read fault to page 'p'
  - Processor 'i' finds owner of page 'p'
  - Owner of 'p' sends copy of page to 'i' and adds 'i' to copyset of 'p'. Processor 'i' has read-only access to 'p'

# Case studies: IVY (Cont.)

- Algorithms used for implementing actions for 'Read' and 'Write' actions
- Centralized manager scheme
  - Central manager resides on single processor: maintains all data ownership information
  - On page fault, processor 'i' requests copy of page from central manager
  - Central manager sends request to page owner. If 'Write' requested, updates owner information to indicate 'i' is the new owner
  - Owner sends copy of page to processor 'i' and
    - If 'Write', also sends copyset of page
    - If 'Read', adds 'i' to the copyset of page
  - On write, central manager sends invalidation messages to all processors in copyset
  - Performance issues
    - Two messages are required to locate page owner
    - On 'Writes', invalidation messages are sent to all processors in copyset
    - Centralized manager can become bottleneck

# Algorithms used for implementing actions for 'Read' and 'Write' actions

- The fixed distributed manager scheme
  - Distributes the central manager's role to every processor in the system
  - Every processor keeps track of the owners of a predetermined set of pages (determined by a mapping function  $H$ )
  - When a processor 'i' faults on page 'p', processor 'i' contacts processor  $H(p)$  for a copy of the page
  - The rest the protocol is the same as the one with the centralized manager

*Note:* In both the centralized and fixed distributed manager schemes, if two or more concurrent accesses to the same page are requested, the requests are serialized by the manager

# Algorithms used for implementing actions for 'Read' and 'Write' actions

- The dynamic distributed manager scheme
  - Every host keeps track of the ownership of the pages that are in its local page table
    - Every page table has a field called *probowner* (probable owner)
    - Initially, *probowner* is set to a default processor
    - The field is modified as pages are requested from various processors
  - When a processor has a page fault, it sends a page request to processor 'i' indicated by the *probowner* field
  - If processor 'i' is the true owner of the page, fault handling proceeds like in centralized scheme
  - If 'i' is not the owner, it forwards the request to the processor indicated in its *probowner* field
  - This continues until the true owner of the page is found

# Case studies: Mirage

- Developed at UCLA, kernel modified to support DSM operation
- Extends the coherence protocol of IVY system to control thrashing (in IVY, a page can move back and forth between multiple processors sharing the page)
- When a shared memory page is transferred to a processor, that processor will keep the page for 'delta' seconds
  - If a request for the page is made before 'delta' seconds expired, processor informs control manager of the amount of time left
  - 'Delta' can be a combination of real-time and service-time for that processor
- Advantages
  - Benefits locality of reference
  - Decreases thrashing

# Case studies: Clouds

- Developed at Georgia Institute of Technology
- The virtual address space of all objects is viewed as a global distributed shared memory
  - The objects are composed of segments which are mapped into virtual memory by the kernel using the memory management hardware
  - A segment is a multiple of the physical page size
- For remote object invocations, the DSM mechanism transfers the required segments to the requesting host
  - On a segment fault, a *location system object* is consulted to locate the object
  - The *location system object* broadcasts a query for each locate operation
  - The actual data transfer is done by the distributed shared memory controller (DSMC)

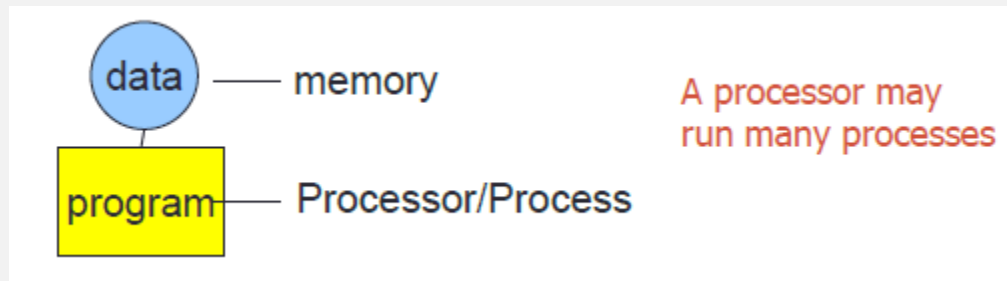
# Message Passing & Programming Using the Message- Passing Paradigm




- Slides Taken from Hanjun Kim,  
Princeton University

# The Message-Passing Programming Paradigm

- Sequential Programming Paradigm



- Message-Passing Programming Paradigm

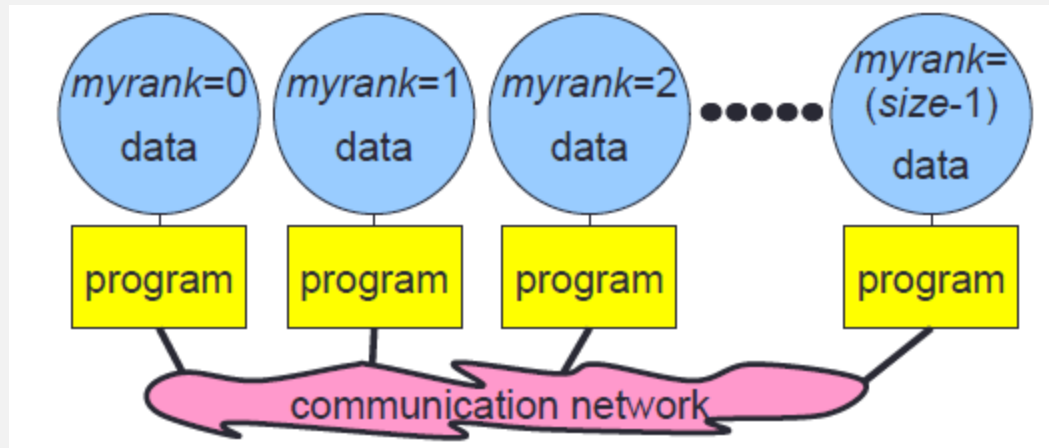


# MPI Operation

- A **process** is a program performing a task on a **processor**
- Each processor/process in a message passing program runs a instance/copy of a program
- Typically a single program operating of multiple dataset
- The variables of each sub-program have
  - The same name
  - But different locations (distributed memory) and different data
  - i.e., all variables are local to a process
- Communicate via special send and receive routines (message passing)

# Data and Work Distribution

- To communicate together MPI processes need identifiers: **rank = identifier number**
- All distribution decision are based on the **rank**
- i.e., which process works on which data



# SPMD example



```
main(int argc, char **argv){  
    if(process is assigned Master role){  
        /* Assign work and coordinate workers  
and collect results */  
        MasterRoutine(/*arguments*/);  
    } else { /* it is worker process */  
        /* interact with master and other  
workers. Do the work and send results to  
the master*/  
        WorkerRoutine(/*arguments*/);  
    }  
}
```

# Why MPI?

- Small:
  - Many programs can be written with only 6 basic functions
- Large:
  - MPI's extensive functionality from many functions
- Scalable:
  - Point-to-point communication
- Flexible:
  - Don't need to rewrite parallel programs across platforms

# Message Passing

- Messages are packets of data moving between sub-programs
- Necessary information for the message passing system:
- Sending Process, Receiving process -->i.e., the ranks
- Source location, Destination Location
- Source Data type, Destination Data type
- Source Data Size, Destination buffer size



# Access

- A sub-program needs to be connected to a message passing system
- A message passing system is similar to:
  - phone line
  - mail box
  - fax machine
  - etc.
- MPI:
  - program must be linked with an MPI library
  - program must be started with the MPI startup tool



# Basic functions

FUNCTION	DESCRIPTION
<code>int MPI_Init(int *argc, char **argv)</code>	Initialize MPI
<code>int MPI_Finalize()</code>	Exit MPI
<code>int MPI_Comm_size(MPI_Comm comm, int *size)</code>	Determine number of processes within a comm
<code>int MPI_Comm_rank(MPI_Comm comm, int *rank)</code>	Determine process rank within a comm
<code>int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)</code>	Send a message
<code>int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int src, int tag, MPI_Comm comm, MPI_Status *status)</code>	Receive a message

# Communicator

- An identifier associated with a group of processes
  - Each process has a **unique rank** within a specific communicator from 0 to (nprocesses-1)
  - Always required when initiating a communication by calling an MPI function
- Default: **MPI\_COMM\_WORLD**
  - contains all processes
- Several communicators can co-exist
  - A process can belong to different communicators at the same time

# Hello World

- `#include "mpi.h"`
- `intmain( intargc, char *argv[] ) {`
- `intnproc, rank;`
- `MPI_Init(&argc,&argv); /* Initialize MPI*/`
- `MPI_Comm_size(MPI_COMM_WORLD,&nproc); /* Get CommSize*/`
- `MPI_Comm_rank(MPI_COMM_WORLD,&rank); /* Get rank*/`
- `printf("Hello World from process %d\n", rank);`
- `MPI_Finalize(); /* Finalize*/`
- `return 0;`
- `}`

# How to compile


- Fortunately, most MPI implementations come with scripts that take care of these issues:
  - `mpicc mpi_code.c -o a.out`
- Two widely used (and free) MPI implementations
  - MPICH (<http://www-unix.mcs.anl.gov/mpi/mpich>)
  - OPENMPI (<http://www.openmpi.org>)

# Point-to-Point Communication

- Simplest form of message passing.
- One process sends a message to another.
- Different types of point-to-point communication:
  - synchronous send
  - buffered = asynchronous send


# Synchronous Sends

- The sender gets an information that the message is received.
- Analogue to the *beep* or *okay-sheet* of a fax.



# Buffered = Asynchronous Sends

- Only know when the message has left.



# Blocking Operations

- **Some sends/receives may block until another process acts:**
- Synchronous send operation **blocks until** receive is issued;
- Receive operation **blocks until** message is sent.
- Blocking subroutine returns only when the operation has completed.



# Blocking Message Passing

- The call waits until the data transfer is done:
- The sending process waits until all data are transferred to the system buffer
- The receiving process waits until all data are transferred from the system buffer to the receive buffer
- Buffers can be freely reused

# Blocking Message Send

- **MPI\_Send(void \*buf, intcount, MPI\_Datatype dtype, int dest, int tag, MPI\_Comm comm);**

• <b>buf</b>	Specifies the starting address of the buffer.
• <b>count</b>	Indicates the number of buffer elements
• <b>dtype</b>	Denotes the datatype of the buffer elements
• <b>dest</b>	Specifies the rank of the destination process in the group associated with the communicator comm
• <b>tag</b>	Denotes the message label
• <b>comm</b>	Designates the communication context that identifies a group of processes

# Blocking Message Send

Standard (MPI_Send)	The sending process returns when the system can buffer the message or when the message is received and <b>the buffer is ready for reuse.</b>
Buffered (MPI_Bsend)	The sending process returns when the message is buffered in <b>an application-supplied buffer.</b>
Synchronous (MPI_Ssend)	The sending process returns only if a matching receive is posted and <b>the receiving process has started to receive the message.</b>
Ready (MPI_Rsend)	The message is <b>sent as soon as possible.</b>

# Blocking Message Receive

- **MPI\_Recv(void \*buf, int count, MPI\_Datatype dtype, int source, int tag, MPI\_Comm comm, MPI\_Status \*status);**


• <b>buf</b>	Specifies the starting address of the buffer.
• <b>count</b>	Indicates the number of buffer elements
• <b>dtype</b>	Denotes the datatype of the buffer elements
• <b>source</b>	Specifies the rank of the source process in the group associated with the communicator comm
• <b>tag</b>	Denotes the message label
• <b>comm</b>	Designates the communication context that identifies a group of processes
• <b>status</b>	Returns information about the received message

# Example

- ...
- `if (rank == 0)`
- `{`
- `for (i=0; i<10; i++) buffer[i] = i;`
- `MPI_Send(buffer, 10, MPI_INT, 1, 123, MPI_COMM_WORLD);`
- `}`
- `else if (rank == 1)`
- `{`
- `for (i=0; i<10; i++)buffer[i] = -1;`
- `MPI_Recv(buffer, 10, MPI_INT, 0, 123, MPI_COMM_WORLD, &status);`
- `for (i=0; i<10; i++)`
- `if (buffer[i] != i)`
- `printf("Error: buffer[%d] = %d but is expected to be %d\n", i, buffer[i], i);`
- `}`
- ...

# Non-Blocking Operations

- Non-blocking operations return immediately and allow the sub-program to perform other work.



# Non-blocking Message Passing

- Returns immediately after the data transferred is initiated
- Allows to overlap computation with communication
- Need to be careful though
  - When send and receive buffers are updated before the transfer is over, the result will be wrong

# Non-blocking Message Passing

- `MPI_Isend(void *buf, int count, MPI_Datatype dtype, int dest, int tag, MPI_Comm comm, MPI_Request *req);`
- `MPI_Recv(void *buf, int count, MPI_Datatype dtype, int source, int tag, MPI_Comm comm, MPI_Request *req);`
- `MPI_Wait(MPI_Request *req, MPI_Status *status);`

• `req`

Specifies the request used by a completion routine when called by the application to complete the send operation.

Blocking	<code>MPI_Send</code>	<code>MPI_Bsend</code>	<code>MPI_Ssend</code>	<code>MPI_Rsend</code>	<code>MPI_Recv</code>
Non-blocking	<code>MPI_Isend</code>	<code>MPI_Ibsend</code>	<code>MPI_Issend</code>	<code>MPI_Irsend</code>	<code>MPI_Irecv</code>



# Non-blocking Message Passing

```
...
right = (rank + 1) % nproc;
left = rank - 1;
if (left < 0)          left = nproc - 1;
MPI_Irecv(buffer, 10, MPI_INT, left, 123,
  MPI_COMM_WORLD, &request);
MPI_Isend(buffer2, 10, MPI_INT, right, 123,
  MPI_COMM_WORLD, &request2);
MPI_Wait(&request, &status);
MPI_Wait(&request2, &status);
...
```

# How to execute MPI codes?

- The implementation supplies scripts to launch the MPI parallel calculation
- **Mpirun -np #proc a.out**
- **Mpiexec -n #proc a.out**
- A copy of the same program runs on each processor core within its own process (private address space)

# Collective Communications


- A single call handles the communication between all the processes in a communicator
- There are 3 types of collective communications
  - Data movement (e.g. MPI\_Bcast)
  - Reduction (e.g. MPI\_Reduce)
  - Synchronization (e.g. MPI\_Barrier)

# Broadcast

- A one-to-many communication.

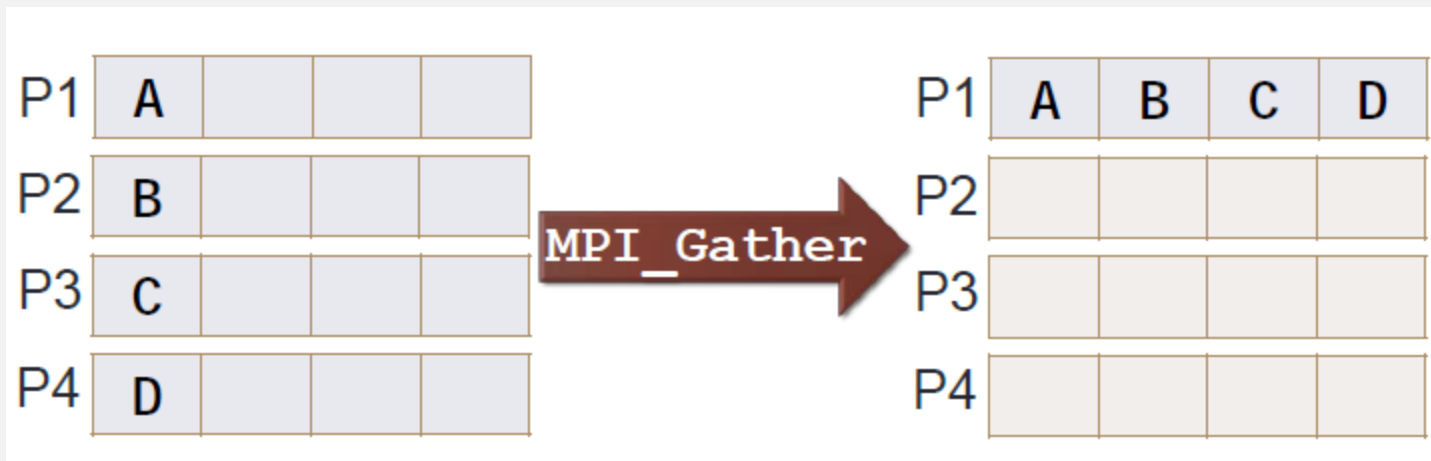


- **Int MPI\_Bcast(void \*buffer, int count, MPI\_Datatype datatype, int root, MPI\_Comm comm);**
- One process (root) sends data to all the other processes in the same communicator
- Must be called by all the processes with the same arguments



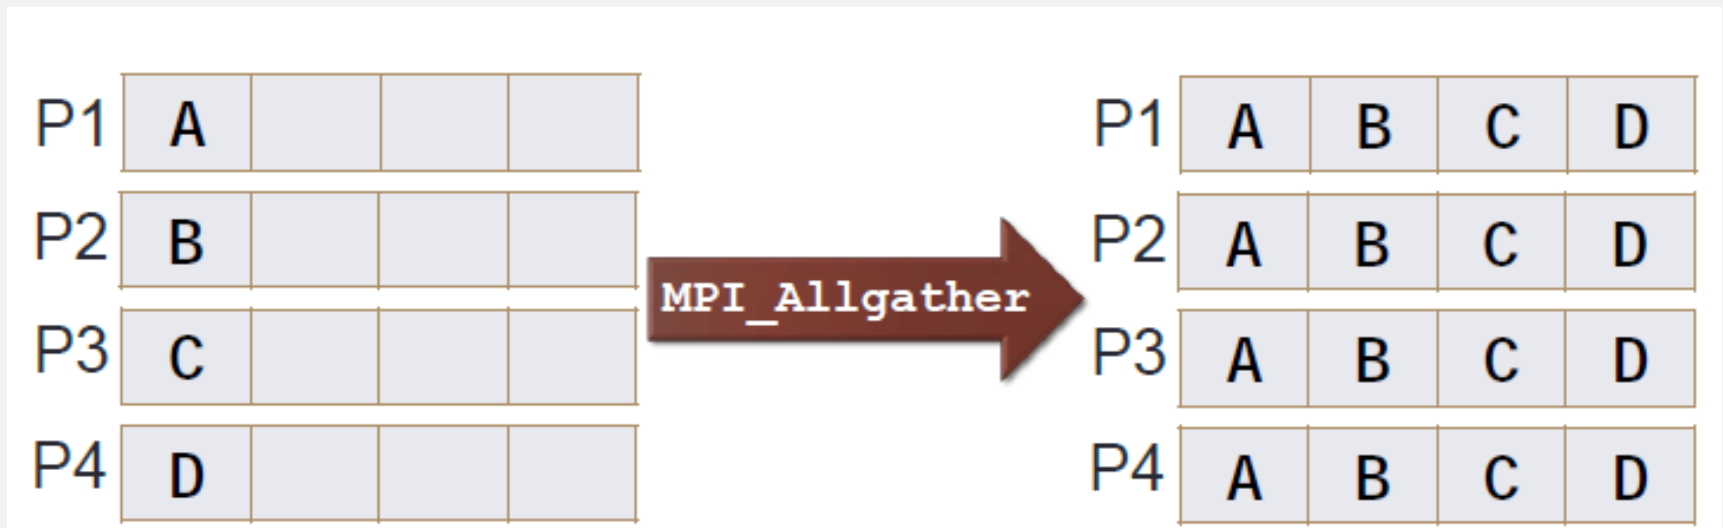
# Gather

- **Int MPI\_Gather(void \*sendbuf, int sendcnt, MPI\_Datatype sendtype, void \*recvbuf, int recvcnt, MPI\_Datatype recvtype, int root, MPI\_Comm comm)**
- One process (root) collects data to all the other processes in the same communicator
- Must be called by all the processes with the same arguments




# Gather to All

- **int MPI\_Allgather(void \*sendbuf, int sendcnt, MPI\_Datatype sendtype, void \*recvbuf, int recvcnt, MPI\_Datatype recvtype, MPI\_Comm comm)**
- All the processes collect data to all the other processes in the same communicator
- Must be called by all the processes with the same arguments



# Reduction Operations

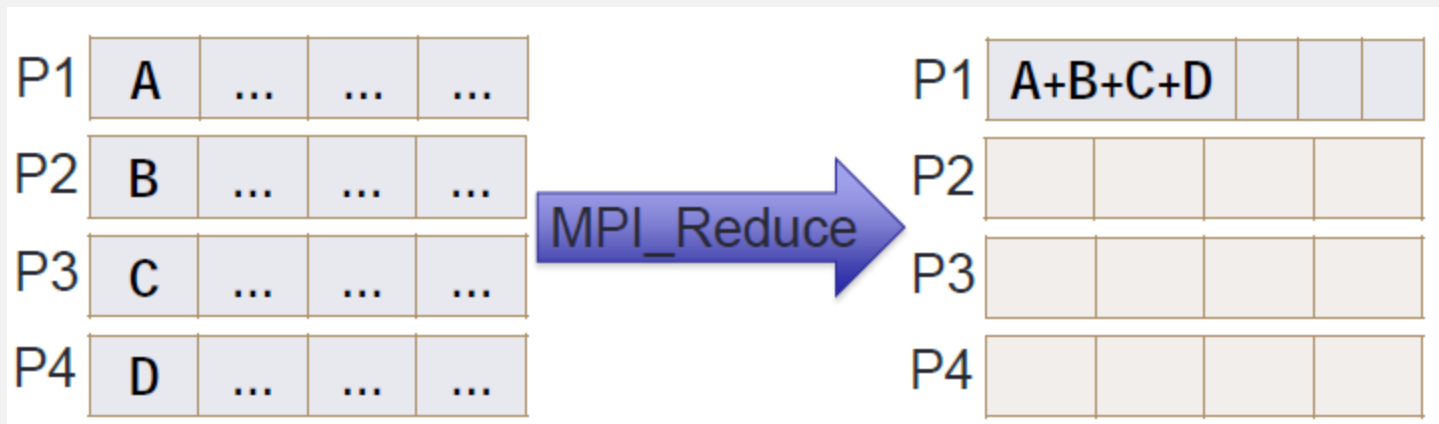
- Combine data from several processes to produce a single result.



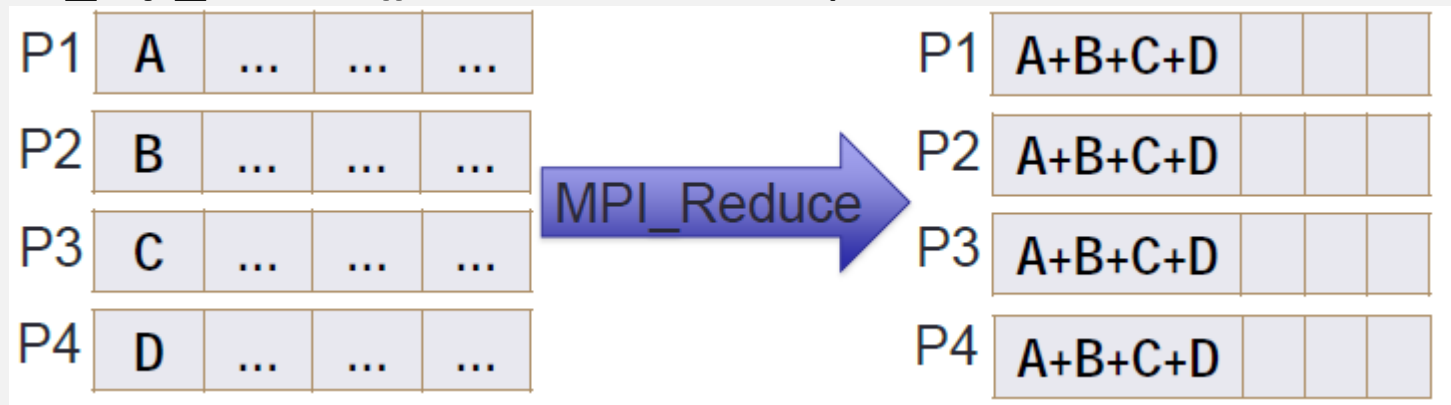


# Reduction

- `int MPI_Reduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)`
- One process (root) collects data to all the other processes in the same communicator, and performs an operation on the data
- `MPI_SUM`, `MPI_MIN`, `MPI_MAX`, `MPI_PROD`, logical AND, OR, XOR, and a few more
- `MPI_Op_create()`: User defined operator




- **int MPI\_Allreduce(void \*sendbuf, void \*recvbuf, int count, MPI\_Datatype datatype, MPI\_Op op, MPI\_Comm comm)**
- All the processes collect data to all the other processes in the same communicator, and perform an operation on the data
- MPI\_SUM, MPI\_MIN, MPI\_MAX, MPI\_PROD, logical AND, OR, XOR, and a few more
- **MPI\_Op\_create():** User defined operator



# Barriers

- Synchronize processes.



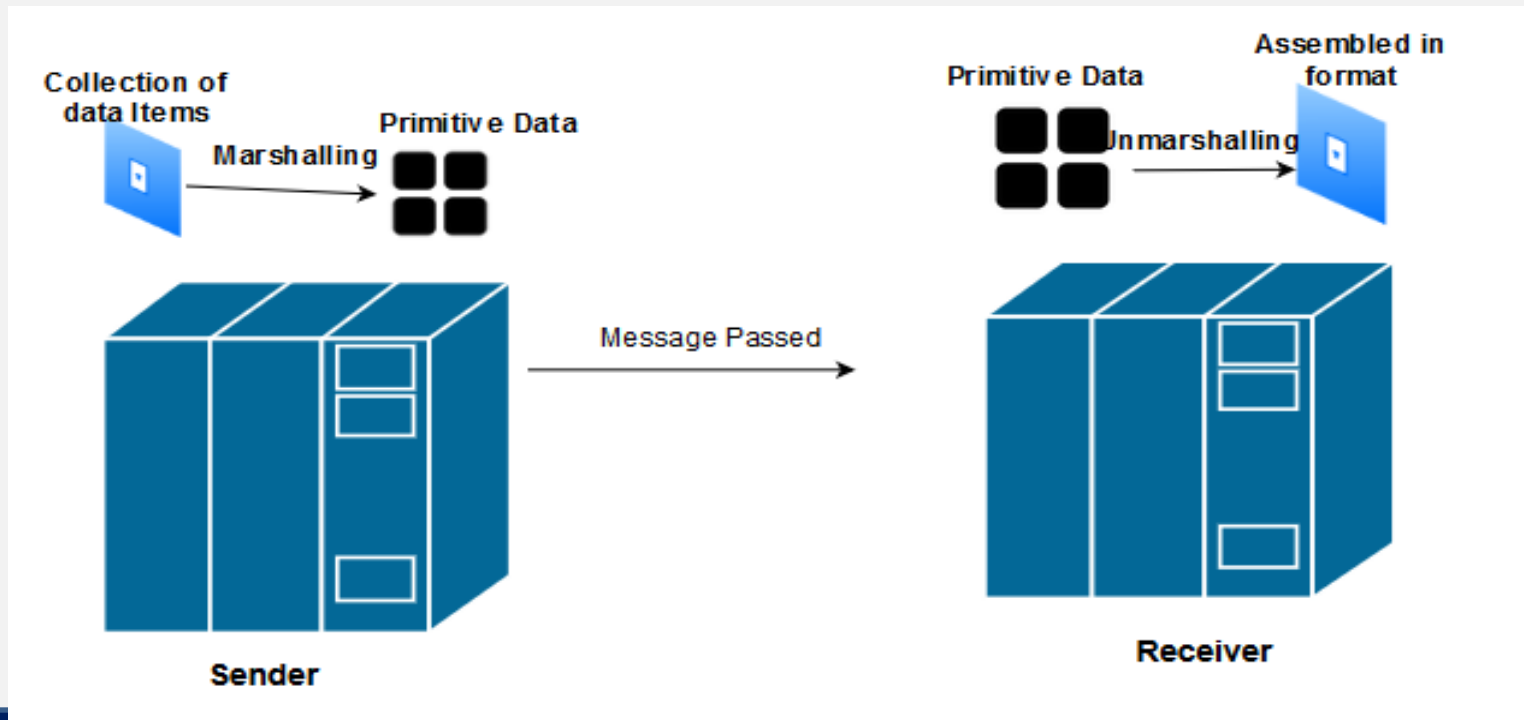
# Synchronization

```
int MPI_Barrier(MPI_Comm comm)

#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[]) {
    int rank, nprocs;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Barrier(MPI_COMM_WORLD);
    printf("Hello, world.  I am %d of %d\n", rank,
nprocs);
    MPI_Finalize();
    return 0;
}
```

# Marshalling

- Process of taking collection of data items and assembling them in order to transmit
- Unmarshalling is the reverse of marshalling




# References

- <http://www.mpi-forum.org>
- <http://www.llnl.gov/computing/tutorials/mpi/>
- <http://www.nersc.gov/nusers/help/tutorials/mpi/intro/>
- <http://www-unix.mcs.anl.gov/mpi/tutorial/gropp/talk.html>
- <http://www-unix.mcs.anl.gov/mpi/tutorial/>
- MPICH (<http://www-unix.mcs.anl.gov/mpi/mpich/>)
- Open MPI (<http://www.open-mpi.org/>)
  
- MPI descriptions and examples are referred from
  - [http://mpi.deino.net/mpi\\_functions/index.htm](http://mpi.deino.net/mpi_functions/index.htm)
  - *Stéphane Ethier (PPPL)'s PICSciE/PICASso Mini-Course Slides*

# Case Study (RPC and Java RMI)

# RPC





# RMI

- The **RMI** (Remote Method Invocation) is an API that provides a mechanism to create distributed application in java.
- The RMI allows an object to invoke methods on an object running in another JVM.
- The RMI provides remote communication between the applications using two objects *stub* and *skeleton*.

# ...Continued...

- A **remote object** is an object whose method can be invoked from another JVM.
- stub
  - The stub is an object, acts as a gateway for the client side.
  - All the outgoing requests are routed through it.
  - It resides at the client side and represents the remote object.

# What happens when the caller invokes method on the stub object?

- It initiates a connection with remote Virtual Machine (JVM),
- It writes and transmits (marshals) the parameters to the remote Virtual Machine (JVM),
- It waits for the result
- It reads (unmarshals) the return value or exception, and
- It finally, returns the value to the caller.


# skeleton

- The skeleton is an object, acts as a gateway for the server side object.
- All the incoming requests are routed through it.

# Upon receiving the incoming request

- It reads the parameter for the remote method
- It invokes the method on the actual remote object, and
- It writes and transmits (marshals) the result to the caller.

# ...Continued...



# Distributed applications

## Requirements


- If any application performs these tasks, it can be distributed application.
  - The application need to locate the remote method
  - It need to provide the communication with the remote objects, and
  - The application need to load the class definitions for the objects.
- The RMI application have all these features, so it is called the distributed application.

# Java RMI Example

- 6 steps to write the RMI program
  - Create the remote interface
  - Provide the implementation of the remote interface
  - Compile the implementation class and create the stub and skeleton objects using the rmic tool
  - Start the registry service by rmiregistry tool
  - Create and start the remote application
  - Create and start the client application



- The client application need only two files, remote interface and client application.
- In the rmi application, both client and server interacts with the remote interface.
- The client application invokes methods on the proxy object, RMI sends the request to the remote JVM.
- The return value is sent back to the proxy object and then to the client application.



# 1) create the remote interface

```
import java.rmi.*;  
public interface Adder extends Remote  
{  
    public int add(int x,int y)throws RemoteException;  
}
```

## 2) Provide the implementation of the remote interface

```
import java.rmi.*;
import java.rmi.server.*;
public class AdderRemote extends UnicastRemote
eObject implements Adder{
AdderRemote()throws RemoteException{
super();
}
public int add(int x,int y){return x+y;}
}
```

### 3) create the stub and skeleton objects using the rmic tool.

- The rmic tool invokes the RMI compiler and creates stub and skeleton objects
- **rmic AdderRemote**

## 4) Start the registry service by the rmiregistry tool

- start the registry service by using the rmiregistry tool.

`rmiregistry 5000`

5000 – port number

# 5) Create and run the server application

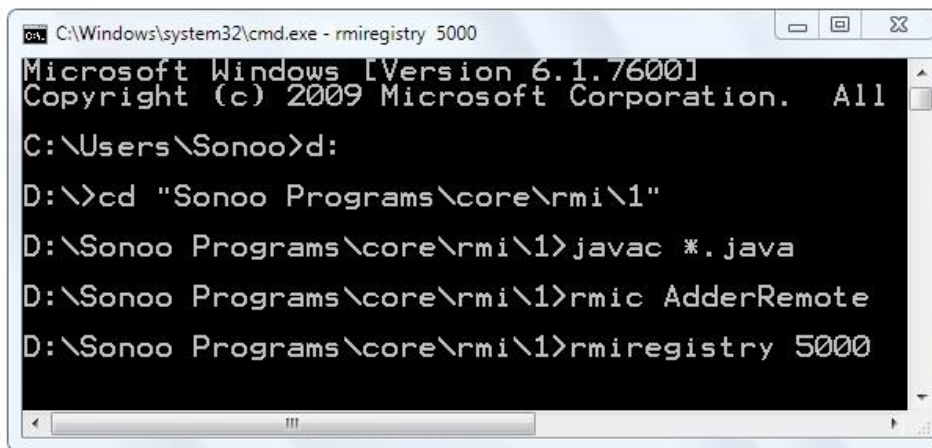
```
import java.rmi.*;
import java.rmi.registry.*;
public class MyServer{
public static void main(String args[]){
try{
Adder stub=new AdderRemote();
Naming.rebind("rmi://localhost:5000/sonoo",stub);
}catch(Exception e){System.out.println(e);}
}
}
```

# 6) Create and run the client application

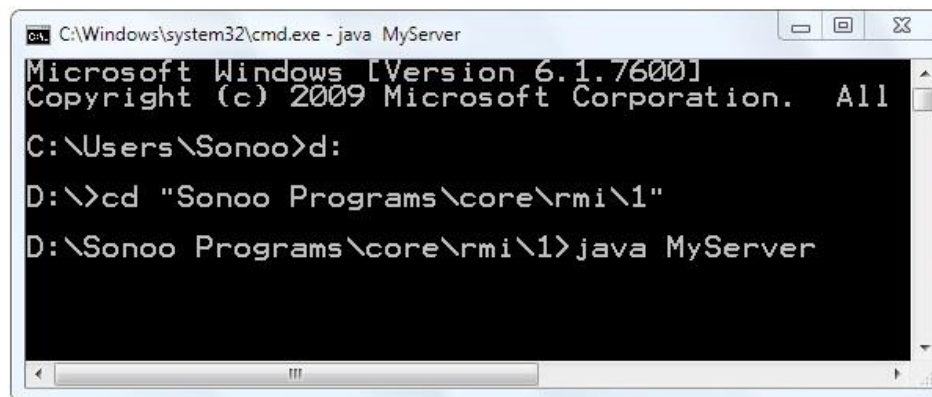
```
import java.rmi.*;
public class MyClient{
public static void main(String args[]){
try{
Adder stub=(Adder)Naming.lookup("rmi://localhost:5000/sonoo");
System.out.println(stub.add(34,4));
}catch(Exception e){}
}
}
```

- For running **this** rmi example,
- 1) compile all the java files : `javac *.java`
- 2) create stub and skeleton object by rmic tool :  
`rmic AdderRemote`
- 3) start rmi registry in one command prompt :  
`rmiregistry 5000`
- 4) start the server in another command prompt :  
`java MyServer`
- 5) start the client application in another command prompt : `java MyClient`





```
C:\Windows\system32\cmd.exe - rmiregistry 5000
Microsoft Windows [Version 6.1.7600]
Copyright (c) 2009 Microsoft Corporation. All
C:\Users\Sonoo>d:
D:\>cd "Sonoo Programs\core\rmi\1"
D:\Sonoo Programs\core\rmi\1>javac *.java
D:\Sonoo Programs\core\rmi\1>rmic AdderRemote
D:\Sonoo Programs\core\rmi\1>rmiregistry 5000
```



```
C:\Windows\system32\cmd.exe - java MyServer
Microsoft Windows [Version 6.1.7600]
Copyright (c) 2009 Microsoft Corporation. All
C:\Users\Sonoo>d:
D:\>cd "Sonoo Programs\core\rmi\1"
D:\Sonoo Programs\core\rmi\1>java MyServer
```



```
C:\Windows\system32\cmd.exe
Microsoft Windows [Version 6.1.7600]
Copyright (c) 2009 Microsoft Corporation. All rights reserved.

C:\Users\Sonoo>d:
D:\>cd "Sonoo Programs\core\rmi\1"
D:\Sonoo Programs\core\rmi\1>java MyClient
38
D:\Sonoo Programs\core\rmi\1>
```